MPI Global-Restart Fault Tolerance Specification
Version 0.1.3
Unofficial, for comment only

Ignacio Laguna and Giorgis Georgakoudis
ilaguna@llnl.gov, georgakoudis1@llnl.gov

Lawrence Livermore National Laboratory

March 9, 2021

Chapter 1

Global-Restart Fault Tolerance

1.1 Introduction

The traditional method to handle process failures in large-scale scientific applications is periodic, global synchronous checkpoint/restart (CPR). When a process failure occurs in a bulk synchronous MPI program, it quickly propagates to other processes so re-starting the application from a previously-saved checkpoint is a simple solution to recover from failures.

A large number of MPI applications already use some form of global synchronous CPR. The goal of global-restart fault tolerance is to provide an easy-to-use interface to improve the efficiency of CPR in bulk synchronous applications by reducing as much as possible the recovery time when failure occurs.

In this chapter, we refer to the global-restart fault tolerance model and interface as the **Reinit** (i.e., re-initialization) model and interface, respectively.
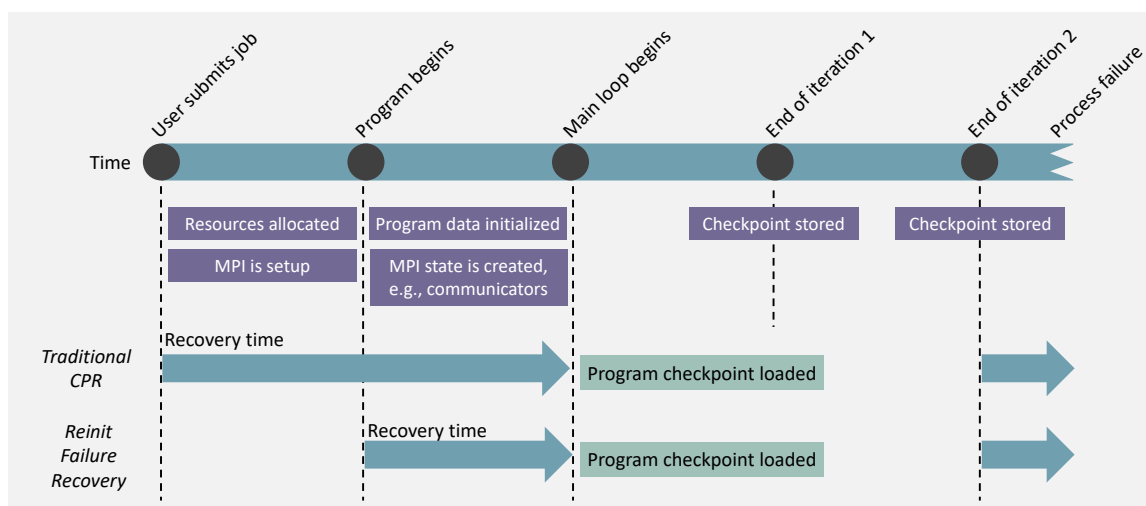

Figure 1.1: The global-restart fault tolerance model (Reinit) provides a mechanisms to reduce the recovery time for bulk synchronous applications that use periodic synchronous checkpoint/restart.

1.2   Fault Model

The Reinit model provides a pre-defined fault-tolerance mechanism to survive **MPI process failures**. We use the definition of process failures used in Section 2.8, i.e., a process failure occurs when an MPI process unexpectedly and permanently stops communicating (e.g., a software or hardware crash results in an MPI process terminating unexpectedly). In the rest of the chapter, when we refer to *failures* we mean *MPI process failures*. The Reinit model assumes that the application's data will be recovered after a failure using a checkpoint that was saved before the failure occurred.

1.3   Reinit MPI Interface

The Reinit interface for global-restart fault tolerance is composed of two MPI functions: `MPI_REINIT` and `MPI_TEST_FAILURE`. This section describes the syntax of these MPI functions.

MPI_Reinit

```
int MPI_Reinit(resilient_fn, void *data)
```

   IN   resilient_fn   user defined procedure (function)
   IN   data   pointer to user defined data

The user-defined procedure should be in C, a function of type MPI_Reinit_function which is defined as:

```
typedef MPI_Reinit_fn void (*)(void *data));
```

The first argument is a user defined procedure, `resilient_fn`, which is called by the MPI_Reinit procedure. The second argument is a pointer to user defined data. This pointer is passed as an argument to the user defined procedure, `resilient_fn`, when the procedure is called. A valid MPI program must contain at most one call to the MPI_Reinit procedure. Calling MPI_Reinit more than one time results in undefined behavior.

The purpose of the user defined `resilient_fn` procedure is to specify a *rollback location*, i.e., a program location to resume execution after a process failure occurs. Depending on the error handler being used, upon the detection of a process failure, MPI will cause the execution of the program to resume at the `resilient_fn` procedure synchronously or asynchronously (see the Error Handling section for more details).

After the `resilient_fn` procedure is re-executed due to failure recovery, the only valid communication objects are the communicators MPI_COMM_WORLD, MPI_COMM_SELF, MPI_COMM_NULL.

> *Advice to users.*   MPI objects that are created before MPI_Reinit is called will not be valid when the `resilient_fn` procedure is re-executed due to a failure. (*End of advice to users.*)

Calling the MPI_Reinit procedure sets the `resilient_fn` procedure to be a rollback location and makes this rollback location active. After activating the rollback location, MPI_Reinit calls the `resilient_fn` procedure. After the MPI_Reinit procedure returns, the rollback location becomes inactive. If a failure occurs during an inactive rollback location, MPI cannot resume execution at the rollback location, and as a result cannot recover from failures using the Reinit model.

Advice to users. To able to survive most of the process failures that can occur during the execution of the program, most calls to MPI and computation should be executed before MPI_Reinit returns. (*End of advice to users.*)

An MPI process must invoke MPI_FINALIZE only after MPI_Reinit returns.

MPI_Test_failure

```
int MPI_Test_failure()
```

The MPI_Test_failure procedure causes the program to resume execution at the rollback point that was activated by MPI_Reinit when two conditions occur: (1) the MPI_ERRORS_REINIT_SYNC handler is associated with MPI_COMM_WORLD, and (2) a failure has been detected before MPI_Test_failure is called.

If no failures were detected before MPI_Test_failure is called, the return code value is MPI_SUCCESS and the procedure performs no operations. If on the other hand failures are detected before the procedure is called, the procedure does not return and it immediately resumes execution at the rollback point.

1.4 Error Handling

MPI provides two predefined error handlers that can be used to handle failures using the Reinit model. These error handlers are intended to be used to handle failures when the World Model is used to initialize MPI. The Reinit error handlers have no effect when the Sessions Model is used.

Unlike other predefined error handlers, such as MPI_ERRORS_ARE_FATAL, that can be associated to communicator, window, file, and session objects, the Reinit error handlers must be associated only to the predefined MPI_COMM_WORLD communicator in the World Model. Associating the Reinit error handlers to window, file, session objects, or communicators other than MPI_COMM_WORLD is undefined.

Rationale. Associating the Reinit error handler to MPI_COMM_SELF would have no effect if a failure occurs because the process that contains MPI_COMM_SELF failed and the error handler cannot be called. Since a process failure during the handling of MPI objects, such as windows, files and sessions eventually manifest itself as a process failure in MPI_COMM_WORLD, it makes sense to associate a Reinit error handler to MPI_COMM_WORLD only. (*End of rationale.*)

The following Reinit error handlers are available in MPI:

- **MPI_ERRORS_REINIT_ASYNC**: The handler is called by MPI immediately after a process failure is detected. The handler, when called, causes the execution of the program to resume at (or jump back to) the active rollback location that was activated by MPI_Reinit.

- **MPI_ERRORS_REINIT_SYNC**: The handler has two effects. The first effect is that it enables the MPI_Test_failure function to cause the execution of the program to resume at (or jump back to) the active rollback location when MPI_Test_failure is called. The second effect is that it returns the error code to the user.

Using the MPI_ERRORS_REINIT_ASYNC handler causes MPI to resume execution of the program when an error is detected whether or not the error is detected during a call to MPI. On the other hand, using the MPI_ERRORS_REINIT_SYNC handler causes MPI to resume execution only after MPI_Test_failure function is called if an error was detected.

1.4.1   Association of Error Handlers

The Reinit error handlers must be associated to MPI_COMM_WORLD before the MPI_-Reinit procedure is called. Calling MPI_Reinit before associating any of the Reinit error handlers produces undefined behavior.

After a Reinit error handler has been associated to MPI_COMM_WORLD, it is invalid to associate a different Reinit error handler to MPI_COMM_WORLD.

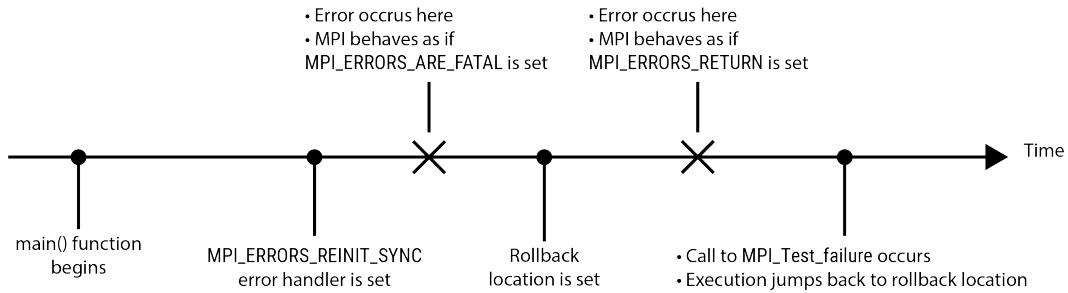Figure 1.2: Different error scenarios for the MPI_ERRORS_REINIT_SYNC error handler.

1.4.2   Behavior for Specific Error Conditions

If an error occurs and one of the Reinit error handlers has been set but there is no active Reinit rollback location, MPI will behave as if the MPI_ERRORS_ARE_FATAL error handler is set (see Figure 1.2).

Errors can occur between the moment the MPI_ERRORS_REINIT_SYNC handler is set and the MPI_Test_failure function is called. If an error occurs in such period of time, MPI behaves as if the MPI_ERRORS_RETURN handler is set.

1.4.3   State

1.5   Examples

Example 1.1   Using Reinit with asynchronous error handling to recover from process failures

```
typedef struct {
    int argc;
    char **argv;
} data_t;

void resilient_function(void *arg)
{
    data_t *data = (data_t *)arg;
    // Cleanup library, if needed
    cleanup_library_state();
    // Resume computation from checkpoint
    // or initialize application data
    if( load_checkpoint() )
        printf("Resume from checkpoint\n");
    else
        init_app_data(data->argc, data->argv);
    bool done = false;
    while(!done) {
        done = compute();
        store_checkpoint();
    }
}

int main(int argc, char *argv[])
{
    // Initialize user defined data type
    data_t data = { argc, argv };

    MPI_Init(argc, argv);
    MPI_Comm_set_errhandler(MPI_COMM_WORLD, MPI_ERRORS_REINIT_ASYNC);
    // MPI_Reinit sets the rollback location
    // to resilient_function and calls it.
    // In asynchronous error handling, the program
    // will go to the rollback location as soon a
    // failure is detected
    MPI_Reinit(&data, resilient_function);
    MPI_Finalize();

    return 0;
}
```

Example 1.2   Using Reinit with synchronous error handling to recover from process failures

```
void resilient_function(void *arg)
{
```

```
1      data_t *data = (data_t *)arg;
2      // Cleanup library, if needed
3      cleanup_library_state();
4      // Resume computation from checkpoint
5      // or initialize application data
6      if( load_checkpoint() )
7          printf("Resume from checkpoint\n");
8      else
9          init_app_data(data->argc, data->argv);
10     bool done = false;
11     while(!done) {
12         done = compute();
13         MPI_Test_failure();
14         store_checkpoint();
15         // Calling MPI_Test_failure will go to the
16         // rollback location, that is resilient_function,
17         // in case of a failure
18         // MPI + computation
19         compute();
20         MPI_Test_failure();
21         // MPI + computation
22         compute();
23         MPI_Test_failure();
24     }
25  }
26
```

1.6   Changes of this Version

1. Added text to specify behavior under the sessions model.

2. Defined that Reinit has a fallback mode of errors_abort, which specifies what happens when one is outside the Reinit function. We mention that outside of the Reinit function the behavior is as if the default handler is set.

3. Question: What happens if you call Reinit before setting the error handler? We handled the case when we are outside the Reinit section. We specify that that we must call the error handler before calling Reinit; otherwise it is undefined behavior.

4. Question: Can you change the error handler from synch to asynch? We specify that we don't support this. You choose a handler and use it in the entire program.

5. We specify that that the only valid way to set the Reinit error handlers is to pass MPI_COMM_WORLD; otherwise it is not a valid program and it should return an error.

6. Question: What happens when we set the error handler, we execute code and a failure occurs, but we didn't call Reinit? We specify that in this case, we the previously set error handler.

7. Added that we assume that the application's state will be recovered using CPR.

8. Modified Example 1.2: (1) put test_failure before C/R; (2) added compute() functions.

9. We define the state of MPI calls when a failure happens before test_failure is called. In this case Reinit behaves as if MPI_ERRORS_RETURN is set so the user is notified of the error but later when test_failure is called the error is recovered.

1.7 To-Do List

1. Define FORTRAN bindings

2. Define what happens with MPI state in tools (e.g., PMPI tools).

3. Why not having multiple rollback locations? Consider supporting multiple protected blocks. Interesting addition, but it will be considered in future work.

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48